

Wine Quality Prediction Using Machine Learning

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sb
5
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import MinMaxScaler
8 from sklearn import metrics
9 from sklearn.svm import SVC
10 from xgboost import XGBClassifier
11 from sklearn.linear_model import LogisticRegression
12
13 import warnings
14 warnings.filterwarnings('ignore')
15
```

```
1 df = pd.read_csv('/content/winequalityN.csv')
2 print(df.head())
3
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	\
0	white	7.0	0.27	0.36		20.7
1	white	6.3	0.30	0.34		1.6
2	white	8.1	0.28	0.40		6.9
3	white	7.2	0.23	0.32		8.5
4	white	7.2	0.23	0.32		8.5

	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	\
0	0.045	45.0	170.0	1.0010	3.00	
1	0.049	14.0	132.0	0.9940	3.30	
2	0.050	30.0	97.0	0.9951	3.26	
3	0.058	47.0	186.0	0.9956	3.19	
4	0.058	47.0	186.0	0.9956	3.19	


	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6

```
1 df.info()
2
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   type                6497 non-null   object
1   fixed acidity       6487 non-null   float64
2   volatile acidity    6489 non-null   float64
3   citric acid         6494 non-null   float64
4   residual sugar      6495 non-null   float64
5   chlorides           6495 non-null   float64
```


```
6  free sulfur dioxide  6497 non-null  float64
7  total sulfur dioxide 6497 non-null  float64
8  density              6497 non-null  float64
9  pH                  6488 non-null  float64
10 sulphates           6493 non-null  float64
11 alcohol             6497 non-null  float64
12 quality             6497 non-null  int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

```
1 df.describe().T
2
```



	count	mean	std	min	
fixed acidity	6487.0	7.216579	1.296750	3.80000	
volatile acidity	6489.0	0.339691	0.164649	0.08000	
citric acid	6494.0	0.318722	0.145265	0.00000	
residual sugar	6495.0	5.444326	4.758125	0.60000	
chlorides	6495.0	0.056042	0.035036	0.00900	
free sulfur dioxide	6497.0	30.525319	17.749400	1.00000	1
total sulfur dioxide	6497.0	115.744574	56.521855	6.00000	7
densitv	6497.0	0.994697	0.002999	0.98711	

```
1 df.isnull().sum()
2
```




type	0
fixed acidity	10
volatile acidity	8
citric acid	3
residual sugar	2
chlorides	2
free sulfur dioxide	0
total sulfur dioxide	0
density	0
pH	9
sulphates	4
alcohol	0
quality	0
dtype: int64	

```
1 for col in df.columns:
2   if df[col].isnull().sum() > 0:
3     df[col] = df[col].fillna(df[col].mean())
4
5 df.isnull().sum().sum()
6
```

0

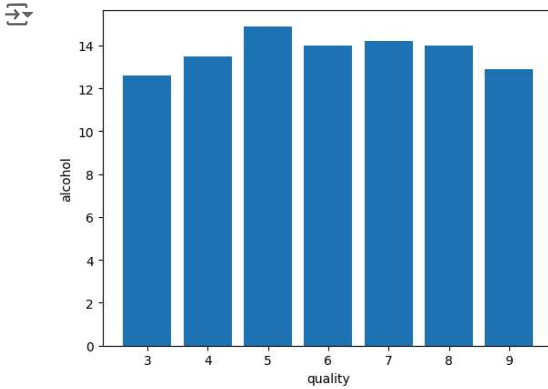
```
1 df.hist(bins=20, figsize=(10, 10))
2 plt.show()
3
```



```

1 plt.bar(df['quality'], df['alcohol'])
2 plt.xlabel('quality')
3 plt.ylabel('alcohol')
4 plt.show()
5

```



```

1 df['best quality'] = [1 if x > 5 else 0 for x in df.quality]
2

```

```

1 df.replace({'white': 1, 'red': 0}, inplace=True)
2

```

```

1 features = df.drop(['quality', 'best quality'], axis=1)
2 target = df['best quality']
3
4 xtrain, xtest, ytrain, ytest = train_test_split(
5     features, target, test_size=0.2, random_state=40)
6
7 xtrain.shape, xtest.shape
8

```

```

↳ ((5197, 11), (1300, 11))

```

```


1 norm = MinMaxScaler()
2 xtrain = norm.fit_transform(xtrain)
3 xtest = norm.transform(xtest)
4

```

```

1 models = [LogisticRegression(), XGBClassifier(), SVC(kernel='rbf')]
2
3 for i in range(3):
4     models[i].fit(xtrain, ytrain)
5
6     print(f'{models[i]} : ')
7     print('Training Accuracy : ', metrics.roc_auc_score(ytrain, models[i].predic
8     print('Validation Accuracy : ', metrics.roc_auc_score(
9         ytest, models[i].predict(xtest)))
10    print()
11

```

 LogisticRegression() :
 Training Accuracy : 0.7019709565048414
 Validation Accuracy : 0.6937888865050418

XGBClassifier(base_score=None, booster=None, callbacks=None,
 colsample_bylevel=None, colsample_bynode=None,
 colsample_bytree=None, device=None, early_stopping_rounds=None,
 enable_categorical=False, eval_metric=None, feature_types=None,
 gamma=None, grow_policy=None, importance_type=None,
 interaction_constraints=None, learning_rate=None, max_bin=None,
 max_cat_threshold=None, max_cat_to_onehot=None,
 max_delta_step=None, max_depth=None, max_leaves=None,
 min_child_weight=None, missing=nan, monotone_constraints=None,
 multi_strategy=None, n_estimators=None, n_jobs=None,
 num_parallel_tree=None, random_state=None, ...) :


Training Accuracy : 0.9735567052182403
 Validation Accuracy : 0.8050515421787681

SVC() :
 Training Accuracy : 0.7069199304892986
 Validation Accuracy : 0.695796426272719

```

1 print(metrics.classification_report(ytest,
2     models[1].predict(xtest)))
3

```



	precision	recall	f1-score	support
0	0.78	0.73	0.75	474
1	0.85	0.88	0.86	826
accuracy			0.83	1300
macro avg	0.81	0.81	0.81	1300
weighted avg	0.82	0.83	0.82	1300

```

1 import matplotlib.pyplot as plt
2 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
3
4 # ... (rest of your code)
5
6 cm = confusion_matrix(ytest, models[1].predict(xtest))
7 disp = ConfusionMatrixDisplay(confusion_matrix=cm)
8 disp.plot()
9 plt.show()

```

